

Devanshu RamaIya

(+91) 9769511290 | devanshu125@gmail.com | linkedin.com/in/devanshu125 | github.com/devanshu125

EDUCATION

NMIMS' MPSTME

Data Science, Bachelor of Technology

Mumbai, India

2019 – 2023

TECHNICAL SKILLS

Languages: Python, SQL, R

ML/AI: PyTorch, TensorFlow, Transformers, Hugging Face, Scikit-learn, LightGBM, XGBoost, LoRA, YOLO

Tools & Infrastructure: Git, Docker, Redis, MongoDB, AWS (SQS), FAISS, Gradio, Streamlit, Pandas, NumPy

EXPERIENCE

El Dorado Capital

Dubai, UAE

Associate (Quantitative Research & Trading)

Apr 2024 – Nov 2025

- Built automated multi-modal AI evaluation system using Gemini 2.0 Flash and GPT-4 to score and rank 500+ entities weekly based on visual pattern recognition and text sentiment analysis, reducing manual review time by 90%
- Designed ML experimentation framework comparing gradient boosting (LightGBM, XGBoost) vs neural network architectures for time-series prediction tasks, processing 50+ GB daily data streams with automated feature engineering and performance benchmarking pipelines
- Engineered high-performance production system with Redis caching for sub-100ms latency on state management and computation, real-time monitoring dashboards, and multi-API orchestration for automated workflow execution

rFund.ai

Remote

Junior Associate (Quantitative Research & Trading)

Dec 2022 – Apr 2024

- Designed and deployed ML pipeline for financial time-series prediction using gradient boosting and traditional ML techniques, implementing automated feature engineering, model training workflows, and rule-based + ML hybrid strategies for systematic signal generation and execution
- Built production-grade multi-broker infrastructure integrating 5+ broker APIs with MongoDB metadata management, AWS SQS for distributed order routing, and Redis-based state management achieving sub-100ms order execution latency
- Engineered real-time monitoring platform with centralized dashboards for position tracking, P&L analytics, margin monitoring, and automated algo control across distributed broker network, supporting \$250K+ in systematic strategy deployment

PROJECTS

Anime Character Attribute Extraction System ↗

- Developed end-to-end pipeline using Vision-Language Models (Qwen-VL) and YOLOv8 for automated character attribute extraction, achieving 96% near-complete extraction rate across 9 attributes with GPU-accelerated preprocessing.
- Engineered scalable architecture reducing per-image processing time from 1.6s to 0.23s (6.9x speedup) on A5000 GPUs through batch processing optimization, and conducted cost-performance analysis for 5M image scaling.
- Implemented 4-bit quantized model variants with comparative evaluation, and built interactive Gradio web application for real-time attribute extraction and visualization.

Reddit Community Rules Violation Classifier ↗

- Engineered multi-model ensemble for rule violation classification by fine-tuning Qwen-4B (LoRA, 4-bit quantization), DeBERTa-v3, and GTE-base models, combining predictions through rank-based normalization to handle model calibration differences.
- Enhanced model performance by implementing semantic search with e5-base model using top-k retrieval and similarity weighted score aggregation.
- Orchestrated parallel training pipeline on dual GPU setup, utilizing data augmentation with positive/negative examples

ACHIEVEMENTS & COMPETITIONS

- Winner, Amazon ML Challenge (NLP) – Rank 7 / 4,898
- Winner, IIIT-D Multilingual Abusive Comment Identification (ShareChat/Moj, NLP) – Rank 1 / 54
- Silver Medal, CommonLit Readability Prediction (Kaggle) – Rank 87 / 3,633
- Finalist, KaggleDays World Championship held in Barcelona sponsored by HP
- Bronze Medal, Optiver Trading at the Close (Kaggle, Time Series) – Rank 283 / 4,436
- Bronze Medal, Jigsaw - Agile Community Rules Classification (Kaggle, LLM) – Rank 157 / 2500
- Bronze Medal, RSNA-MICCAI Brain Tumor Radiogenomic Classification (Kaggle, CV) – Rank 105 / 1555